



KUBERNETES

RESOURCE MANAGEMENT

Christian Dussol



THE RESOURCE MANAGEMENT CHALLENGE

Without governance

- ✗ Pods can consume unlimited resources
- ✗ One misconfigured deployment affects entire cluster
- ✗ Teams exceed budgets without visibility
- ✗ In-place resizing becomes dangerous

THE RESOURCE MANAGEMENT TRIFECTA

Pod Resources

What each container needs/uses
→ Defined in YAML specs

LimitRange

Rules for individual containers
→ Namespace-level constraints

ResourceQuota

Budget for entire namespace
→ Total resource limits



POD RESOURCES

REQUESTS & LIMITS

Requests

What I need

Scheduler allocation

Guaranteed resources

QoS baseline

Limits

Maximum I can use

Protection boundary

Throttling trigger

OOM threshold

POD RESOURCES IN ACTION

```
apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  containers:
  - name: app
    image: nginx
    resources:
```

```
  requests:
```

```
    cpu: "500m"
    memory: "512Mi"
```

```
  limits:
```

```
    cpu: "1"
    memory: "1Gi"
```

What I need to start

Maximum I can
consume



LIMIT RANGE

INDIVIDUAL CONTAINER GUARDRAILS

What it controls

Max size per container

Min size per container

Default values (if not specified)

Why it matters

Prevents oversized containers from destabilizing nodes

Key insight

Validates EVERY pod creation and resize operation

LIMIT RANGE CONFIGURATION

```
apiVersion: v1
kind: LimitRange
metadata:
  name: container-limits
  namespace: my-namespace
spec:
  limits:
  - type: Container
    max:
```

cpu: "4"
memory: "8Gi"

No container can
exceed this

min:

cpu: "100m"
memory: "128Mi"

No container can
be smaller

default:

cpu: "1"
memory: "1Gi"

Applied if not
specified



RESOURCE QUOTA

NAMESPACE BUDGET CONTROL

What it limits

Total CPU across all pods

Total Memory across namespace

Number of pods, services, PVCs

Storage capacity

Why it matters

Prevents teams from consuming cluster resources uncontrollably

Key insight

Validates cumulative usage, not individual pods

RESOURCEQUOTA CONFIGURATION

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: team-quota
  namespace: my-namespace
spec:
  hard:
```

```
requests.cpu: "50"
requests.memory: "100Gi"
```

Total baseline
allocation

```
limits.cpu: "100"
limits.memory: "200Gi"
pods: "100"
```

Total burst capacity

HOW THEY WORK TOGETHER

Developer creates or resizes pod



LimitRange validates
“Is this container size OK?”



ResourceQuota validates
“Does team have budget?”



Kubernetes executes
Pod created or resized safely

The Validation Flow

4

POLICY AS CODE

KYVERNO

Native K8s provides

LimitRange (static rules)

ResourceQuota (hard limits)

Manual enforcement

Kyverno adds

Custom validation rules

Auto-remediation

Advanced conditions

Ratio enforcement

Example

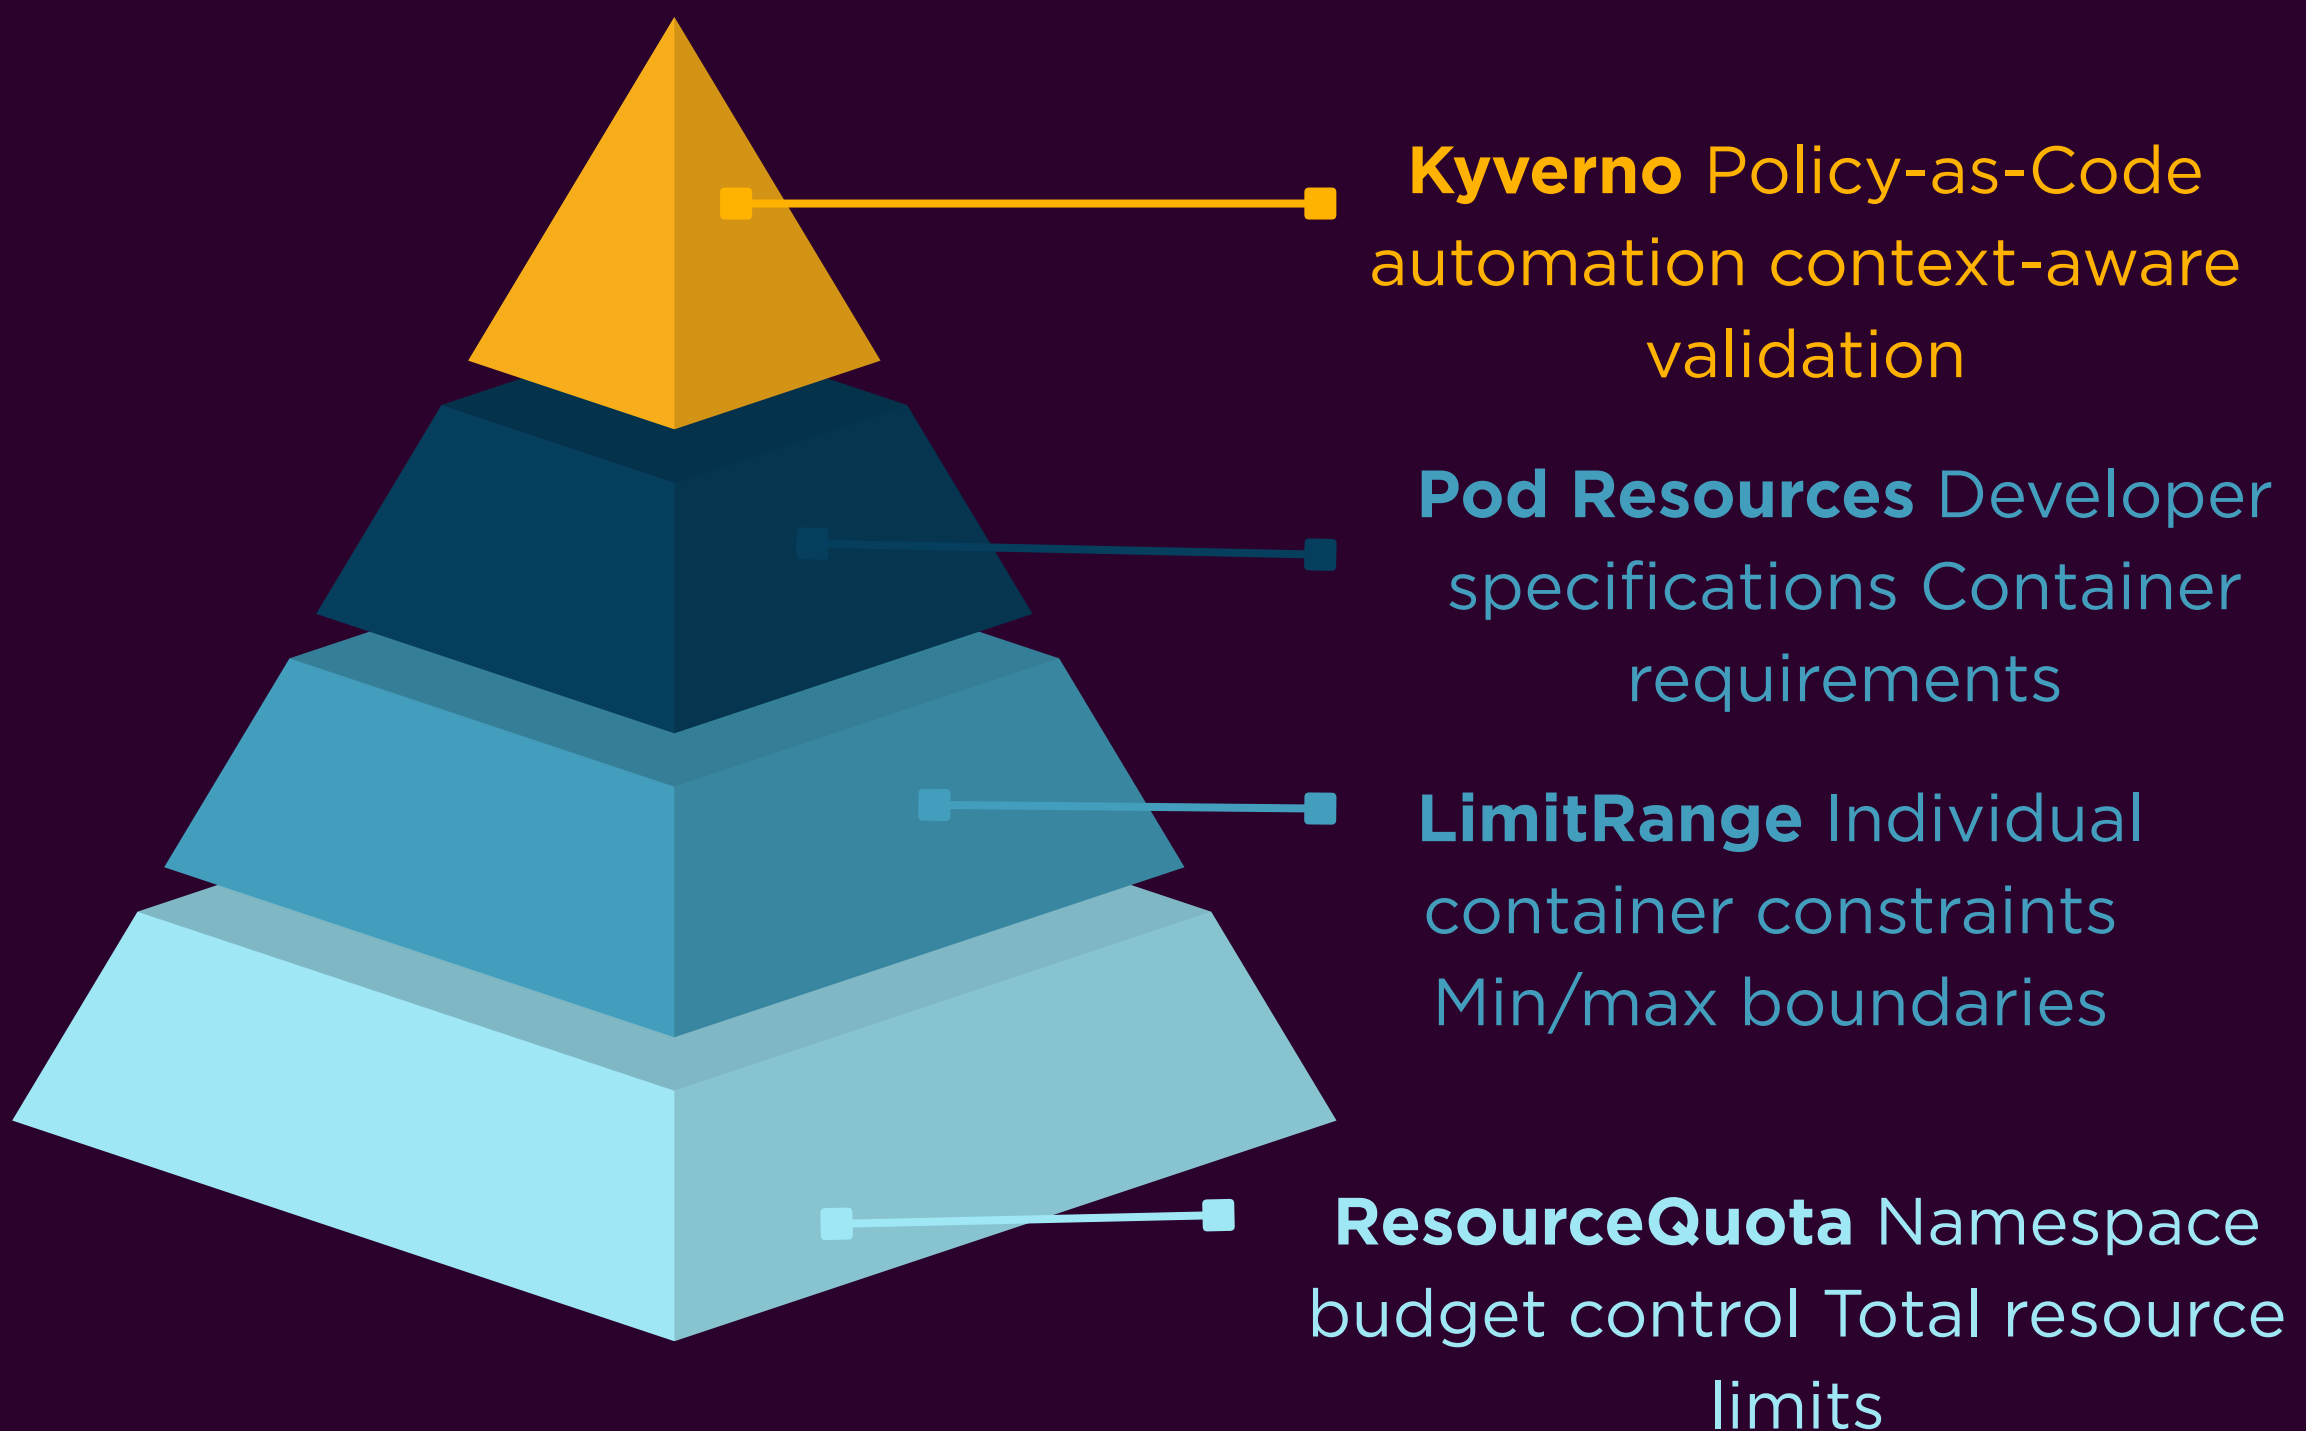
Automatically add resource specs to pods that do not have

KYVERNO AUTO-REMEDiation

```
apiVersion: kyverno.io/v1
kind: ClusterPolicy
metadata:
  name: add-default-resources
spec:
  rules:
  - name: add-resources
    match:
      resources:
        kinds: [Pod]
    mutate:
      patchStrategicMerge:
        spec:
          containers:
            - (name): "*"
              resources:
                +(requests):
                  +(cpu): "100m"
                  +(memory): "128Mi"
```

Automatically adds
resources to pods
→ No manual
intervention
needed

4-LAYER GOVERNANCE



Complete governance + automation

WHY THIS MATTERS FOR POD RESIZING

Without governance

Resize operations can impact other workloads

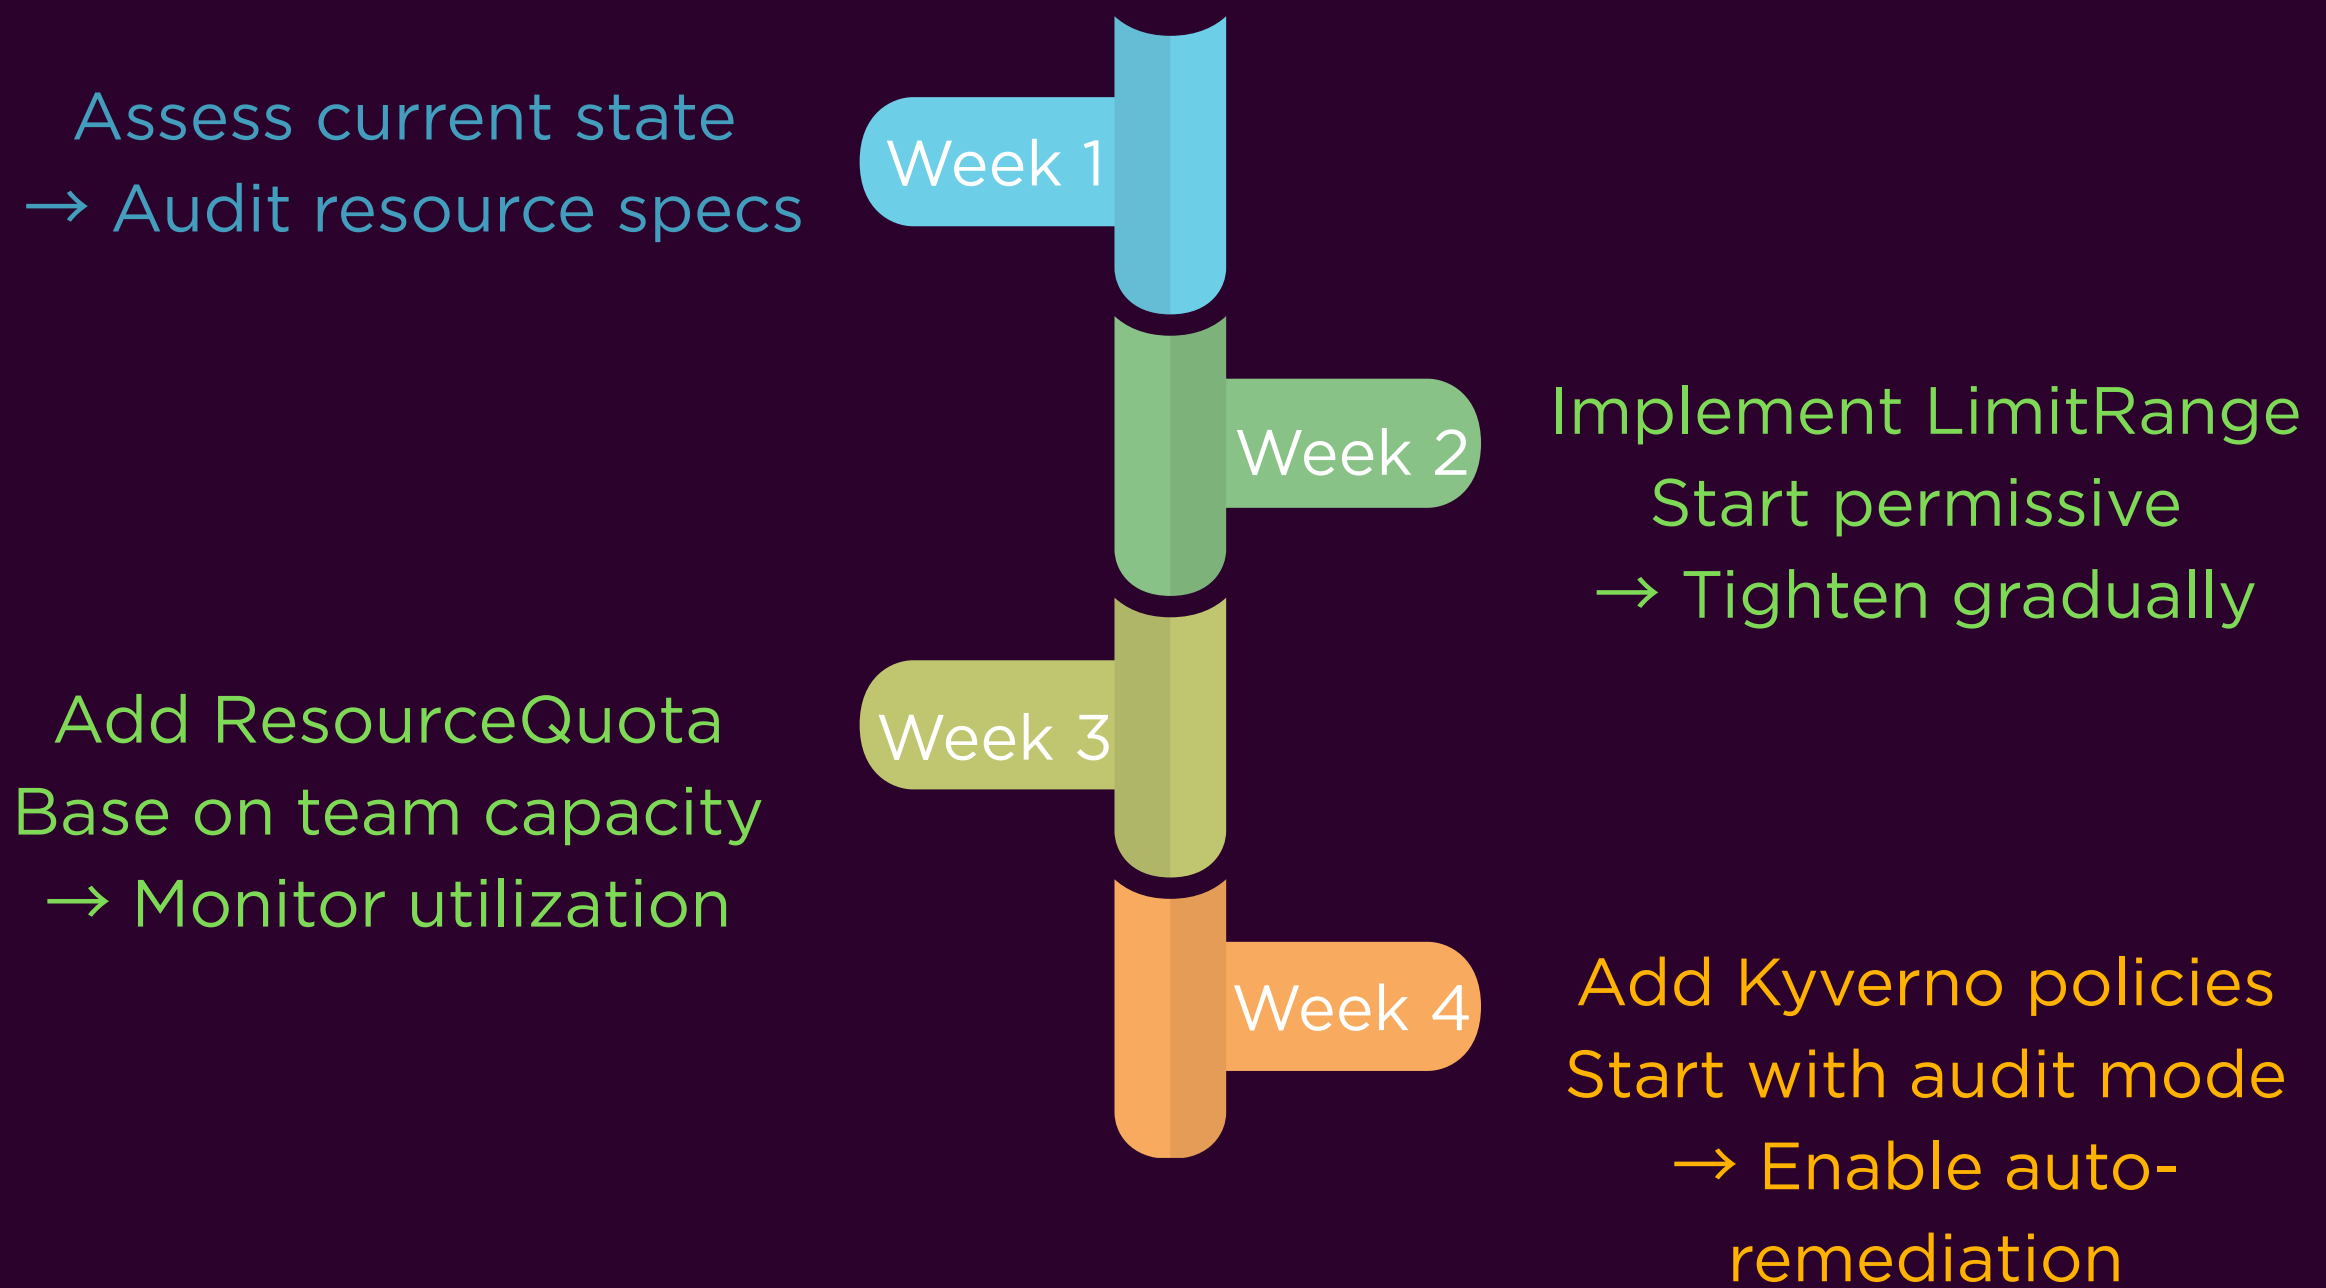
With governance

LimitRange prevents dangerous resizes

ResourceQuota prevents budget overflow

Safe boundaries for dynamic pod sizing

4-WEEK IMPLEMENTATION ROADMAP





Get the Complete Guide

This guide includes:

- ✓ Native K8s resource management (Pod Resources, LimitRange, ResourceQuota)
- ✓ 4-layer governance architecture with Kyverno
- ✓ Production-ready Kyverno policies
- ✓ Complete troubleshooting section with debug commands
- ✓ Best practices for enterprise environments

 GitHub Markdown: <https://bit.ly/490bX2L>

